

Scalability of Distributed Version Control Systems

Michael J. Murphy John Markus Bjørndalen Otto J. Anshus

November 2017

Abstract

Distributed version control is an interesting form of distributed system because it takes eventual consistency to the extreme. Every replica of a repository contains the full history in an append-only data structure, any replica may add new commits, and conflicting updates are reconciled later in a merge operation. These systems are popular, but their use is generally limited to the small text files of source code.

This paper explores the challenges of using version control to store larger binary files, with the goal of building a scalable, highly-available, distributed storage system for media files such as images, audio, and video. We developed an early prototype of such a system, which we call Distributed Media Versioning (DMV). We perform experiments with the popular version control systems Git and Mercurial, the Git-based backup tool Bup, and our DMV prototype.

We measured commit times and repository sizes when storing single files of increasing size, and when storing increasing numbers of single-kilobyte files. We find that processing files whole will limit maximum file size to what can fit in RAM. And we find that storing millions of objects loose as files with hash-based names will result in inefficient write speeds and use of disk space. We conclude that the key to storing large files is to break them into smaller chunks, and that the key to storing many small chunks is to aggregate them into larger files. We intend to incorporate these insights into future versions of DMV.

Incorporate abstract suggestion from Otto

=====Otto=====

A distributed version control system allows for replicas of a repository with the full history of updates. Updates and new commits can be done to each replica independent of each other. Updates, including conflicting updates, are reconciled later in a peer-to-peer merge operation. This will bring replicas to a consistent state with each other.

These systems are popular, but their INTENDED? use is TYPICALLY? (generally) limited to small text files of source code. However, distributed version control can be useful also for larger files, including binary files of images, audio and video.

This paper was presented at the NIK-2017 conference; see <http://www.nik.no/>.

We report on the results from a set of experiments designed to characterise the behaviour of some widely used distributed version control systems. The experiments measured commit times and repository sizes when storing single files of increasing size, and when storing increasing numbers of single-kilobyte files.

The goal is to better understand the behaviour and approaches used in those systems with the purpose of building a scalable, highly-available, distributed storage system for many large text and binary files residing on computers ranging from smart phones to servers.

An early prototype of such a system, Distributed Media Versioning (DMV), is briefly described and compared with Git and Mercurial, and the Git-based backup tool Bup.

We find that processing LARGE files WITHOUT SPLITTING THEM INTO SMALLER PARTS (whole) will limit maximum file size to what can fit in RAM. Storing millions of small files with hash-based names will result in SLOW? (inefficient) write SPEEDS? AS IN ... ms? HIGH LATENCY WRITE OPERATIONS? (speeds) and INEFFICIENT use of disk space. We conclude that large files must be split into smaller chunks to avoid the physical memory limitation, and that storing many small chunks should be aggregated into larger files to achieve low latency write operations and efficient use of disk STORAGE? space.

We intend to incorporate these insights into future versions of DMV.

1 Introduction

Write: Distributed version control is an interesting form of distributed system because it takes eventual consistency to the extreme.

Write: Every replica of a repository contains the full history in an append-only data structure, any replica may add new commits, and conflicting updates are reconciled later in a merge operation.

Written: These systems are popular, but their use is generally limited to the small text files of source code.

Distributed version control systems (DVCSs) are designed primarily to store program source code: plain text files in the range of tens of kilobytes. Checking in larger binary files such as images, sound, or video affects performance. Actions that require copying data in and out of the system slow from hundredths of a second to full seconds or minutes. And since a DVCS keeps every version of every file in every *repository*, forever, the disk space needs compound.

This has lead to a conventional wisdom that binary files should never be stored in version control, inspiring blog posts with titles such as "Don't ever commit binary files to Git! Or what to do if you do" [24], even as the modern software development practice of continuous delivery was commanding teams to "keep absolutely everything in version control." [10, p.33]

Write: This paper explores the challenges of using version control to store larger binary files, with the goal of building a scalable, highly-available, distributed storage system for media files such as images, audio, and video.

2 Distributed Media Versioning

Written: We developed an early prototype of such a system, which we call Distributed Media Versioning (DMV).

Reduce and compress the DMV section

Git's *directed acyclic graph (DAG)* data structure has several interesting properties for distributed data storage.

De-duplication Identical objects are de-duplicated because they will have the same ID and naturally collapse into a single object in the data store. This results in a natural compression of redundant objects.

A record of causality Copies of the DAG can be distributed and updated independently. Concurrent updates will result in multiple branches of history, but references from child commit to parent commit establish a happens-before relationship and give a chain of causality. Branches can be merged by manually reconciling the changes, and then creating a merge commit that refers to both parent commits. When transferring updates from one copy to another, only new objects need to be transferred.

Atomic updates When a new commit is added, all objects are added the database first, then finally the reference to the current commit is updated. This reference is a 160-byte SHA-1 hash value, and which can be updated atomically.

Verifiability Because every object is identified by its cryptographic hash, the data integrity of each object can be verified at any time by re-computing and checking its hash. And because objects refer to other objects by hash, the graph is a form of blockchain. All objects can be verified by checking hashes and following references from the most recent commits down to the initial commits and the blob leaves of the graph.

The efficiency of de-duplication depends on how well identical pieces of data map to identical objects. In Git, the redundant objects are the files and directories that do not change between commits. De-duplication of redundant data within files is accomplished by aggregating objects together into pack files and compressing them with zlib [4, Section 10.4].

Calculating deltas during this *packing* phase requires loading the objects into memory, and so it can cause an out-of-memory error if an object is too large to fit into available RAM. Because Git stores files whole in *binary large objects (blobs)*, it cannot pack objects that are larger than available RAM.

If the DAG operated at a granularity smaller than the file, it could become even more powerful. It could naturally de-duplicate chunks of files the way that Git already de-duplicates whole files, and it could ensure that all objects fit into RAM for packing or other operations.

This sub-file granularity and de-duplication is the core idea behind our new data storage system, *Distributed Media Versioning*.

Distributed Media Versioning (DMV) is our new distributed data storage platform. The core idea is relatively simple — store data in a Git-like DAG, but make the following changes:

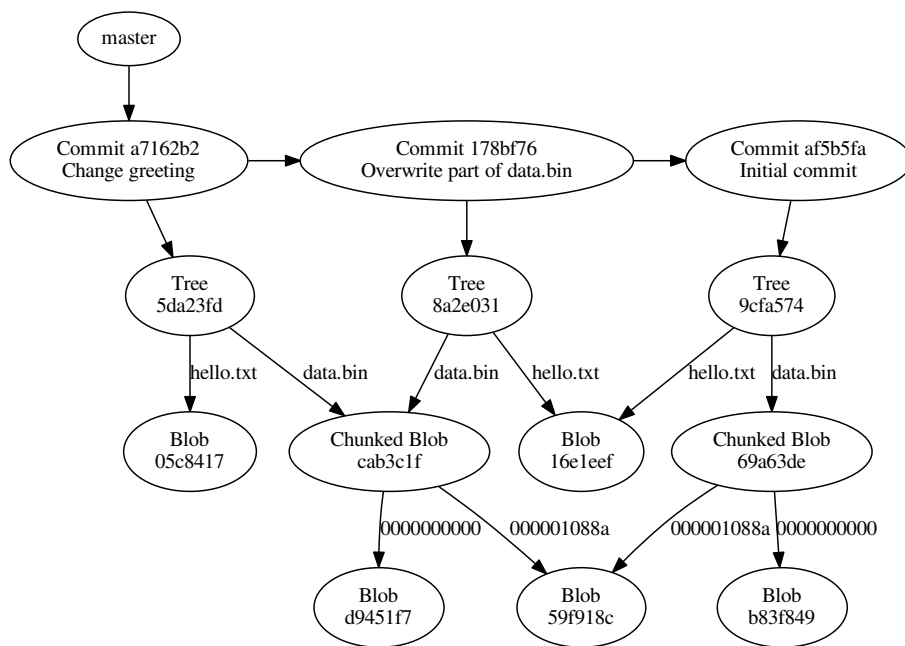


Figure 1: A simple DMV DAG with three commits

1. Store data at a finer granularity than the file
2. Allow nodes to store only a portion of the DAG as a whole

Doing so allows a data set to be replicated or sharded across many nodes according to the capacity of nodes and the needs of local users. The focus is on data locality: tracking what data is where, presenting that information to the user, and making it easy to transfer data to other nodes as desired. The ultimate goal is to create a new abstraction, of *many devices, one data item* in varying states of synchronization.

DMV's data set and its history are represented as a DAG (directed acyclic graph, see ??), and the rest of the design flows from that. DMV's DAG is based on Git's, but it adds a new object type, the *chunked blob*, which represents a blob that has been broken into several smaller chunks. An example DMV DAG is shown in Figure 1.

Files are split into chunks using a *rolling hash algorithm* such as Rabin-Karp fingerprinting [12]. This splits the files into chunks by content rather than position, so that identical chunks within files (and especially different versions of the same file) will be found and stored as identical objects, regardless of their position within the file. This way, identical chunks will be naturally de-duplicated by the DAG, and only the changed portions of files need to be stored as new objects.

It's the
rsync
algorithm

We have written a DMV prototype as a proof-of-concept. The DMV prototype is written in the Rust programming language and it runs from the Unix command line, with the executable built as a thin wrapper around a library, so that it can be used by other applications.

The DMV prototype was developed with Rust stable versions 1.15 and 1.16 on Debian Linux 8.6 ("Jessie"). The current DMV prototype stands at 7592 lines of Rust code (6565 excluding comments). Source code is available alongside this dissertation in Munin, the University of Tromsø's open research archive (<http://munin.uit.no>), and continued

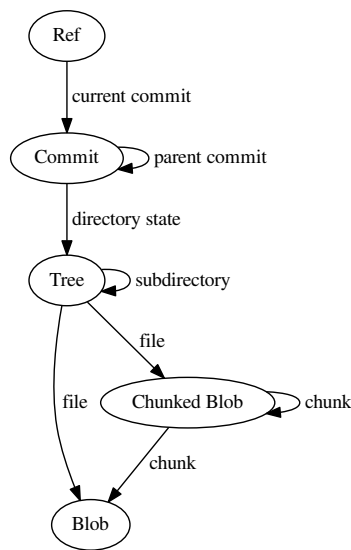


Figure 2: DMV DAG Object Types

work can be found via the author’s website (<http://dmv.sleepymurph.com/>).

3 Experiments

Otto: Don’t use ”would” so much

Written: We perform experiments with the popular version control systems Git and Mercurial, the Git-based backup tool Bup, and our DMV prototype.

Written: We measured commit times and repository sizes when storing single files of increasing size, and when storing increasing numbers of single-kilobyte files.

Methodology

We conducted two major experiments. In order to measure the effect of file size, we would *commit* a single file of increasing size to each target *version control system* (VCS). And to measure the effect of numbers of files, we would commit increasing number of small (1 KiB) files to each target VCS.

We ran each experiment with four different VCSs: Git, Mercurial, Bup, the DMV prototype. We chose Git because it is the most popular DVCS in use today [1]. We chose the Mercurial and Bup because they are both related to Git but each store data differently. Git and DMV both store objects in an *object store* directory as a file named for its hash ID. Git has a separate garbage collection step that takes object files and aggregates them into pack files [4, Section 10.7]. Mercurial stores revisions of each file as a base revision followed by a series of deltas [15, Chapter 4], much like previous VCSs such as RCS, CVS, and Subversion [20]. Bup uses Git’s exact data model and pack file format, but Bup breaks files into chunks using a rolling hash, reusing Git’s *tree* object as a chunked blob index¹. Unlike Git, Bup writes to the pack file format directly, without Git’s separate commit and pack steps, and without bothering to calculate deltas [17]. As a control, we

also ran the experiments with a dummy VCS that simply copied the files to a hidden directory.

For each experiment, the procedure for a single trial was as follows:

1. Create an empty repository of the target VCS in a temporary directory
2. Generate target data to store, either a single file of the target size, or the target number of 1 KiB files
3. Commit the target data to the repository, measuring wall-clock time to commit
4. Verify that the first commit exists in the repository, and if there was any kind of error, run the repository's integrity check operation
5. Measure the total repository size
6. Overwrite a fraction of each target file
7. (Number-of-files experiment only) Run the VCS's status command that lists what files have changed, and measure the wall-clock time that it takes to complete
8. Commit again, measuring wall-clock time to commit
9. Verify that the second commit exists in the repository, and if there was any kind of error, run the repository's integrity check operation
10. Measure the total repository size again
11. (File-size experiment only) Run Git's garbage collector (`git fsck`) to pack objects, then measure total repository size again
12. Delete temporary directory and all trial files

We increased file sizes exponentially by powers of two from 1 B up to 128 GiB, adding an additional step at 1.5 times the base size at each order of magnitude. For example, starting at 1 MiB, we would run trials with 1 MiB, 1.5 MiB, 2 MiB, 3 MiB, 4 MiB, 6 MiB, 8 MiB, 12 MiB, and so on.

We increased numbers of files exponentially by powers of ten from one file to ten million files, adding additional steps at 2.5, 5, and 7.5 times the base number at each order of magnitude. For example, starting at 100 files we would run trials with 100, 250, 500, 750, 1000, 2500, 5000, 7500, 10000, and so on.

Input data files consisted of pseudorandom bytes taken from the operating system's pseudorandom number generator (`/dev/urandom` on Linux).

When updating data files for the second commit, we would overwrite a single contiguous section of each file with new pseudorandom bytes. We would start one-quarter of the way into the file, and overwrite 1/1024th of the file's size (or 1 byte if the file was smaller than 1024 KiB). So a 1 MiB file would have 1 KiB overwritten, a 1 GiB file would have 1 MiB overwritten, and so on.

Experiment Platform

We ran the trials on four dedicated computers with no other load. Each was a typical office desktop with a 3.16 GHz 64-bit dual-core processor and 8 GiB of RAM, running Debian version 8.6 ("Jessie"). Each computer had one normal SATA hard disk (spinning platter, not solid-state), and trials were conducted on a dedicated 197 GiB LVM partition formatted with the ext4 filesystem. All came from the same manufacturer with the same specifications and were, for practical purposes, identical.

¹Git can read a repository written by Bup, but it will see the large file as a directory full of smaller chunk files.

Include
platform
table?

We ran every trial four times, once on each of the experiment computers, and took the mean and standard deviation of each time and disk space measurement. However, because the experiment computers are practically identical, there was little variation.

4 Results

Otto: Don't use "would" so much

File Size

File Size Limits: RAM, Time, Disk Space

Written: We find that processing files whole will limit maximum file size to what can fit in RAM.

In the experiments, both Git and Mercurial had file size limits that were related to the size of RAM. Mercurial would refuse to commit a file 2 GiB or larger. It would exit with an error code and print an error message saying "up to 6442 MB of RAM may be required to manage this file." The commit would not be stored, and the repository would be left unchanged. Mercurial needs to be able to fit the file into memory three times over in order to commit it.

Git's commit operation would appear to fail with files 12 GiB and larger. It would exit with an error code and print an error message saying "fatal: Out of memory, malloc failed (tried to allocate 12884901889 bytes)." However, the commit would be written to the repository, and git's fsck operation would report no errors. So the commit operation completes successfully, even though an error is reported.

With files 24 GiB and larger, Git's fsck operation itself would fail. The fsck command would exit with an error code and give a similar "fatal ... malloc" error. However, the file could still be checked out from the repository without error. So we continued the trials assuming that these were also false alarms.

The DMV prototype was able to store a file up to 64 GiB in size, but time became a limiting factor as file size increased. At 96 GiB, our experiment script timed out and terminated the commit after five and a half hours.

Our experiment environment itself limited the largest file stored by any VCS to 96 GiB. Any larger and it was simply impossible to store a second copy of the file on our 197 GiB test partition. Bup was able to store a 96 GiB file with no errors in just under two hours. Git could also store such a large file, but one must ignore the false-alarm "fatal" errors being reported by the user interface.

These findings are summarized in Table 1 and Table 2.

Commit Times for Increasing File Sizes

Written: Give file size result timing

Figure 3 shows the wall-clock time required for the initial commit, adding a single file of the given size to a fresh repository. Over all, the trend is clear and unsurprising: commit time increases with file size. It increases linearly for Git, Mercurial, and Bup. DMV's commit times increase in a more parabolic fashion, which is most apparent in Figure 3e.

Otto: Give reason Mercurial needs 3x RAM

Ask Otto: Should that go in the discussion section?

Otto: Is this A, D, or I

Otto: Why timed out?

Otto: Make clearer that this was just experiment set up and can be changed

Otto: Why parabolic

Ask Otto: In discussion?

Table 1: Observations as file size increases

Size	Observation
1.5 GiB	Largest successful commit with Mercurial
2 GiB	Mercurial commit rejected
8 GiB	Largest successful commit with Git
12 GiB	Git false-alarm errors begin, but commit still intact
16 GiB	Largest successful Git fsck command
24 GiB	Git false-alarm errors begin during fsck, but commit still intact
64 GiB	Largest successful DMV commit
96 GiB	DMV timeout after 5.5 h
96 GiB	Last successful commit with Bup (and Git, ignoring false-alarm errors)
128 GiB	All fail due to size of test partition

Table 2: Effective size limits for VCSs evaluated

Remove this table?

VCS	Effective limit
Git	Commit intact at all sizes, UI reports errors at 12 GiB and larger
Mercurial	Commit rejected at 2 GiB and larger
Bup	Successful commits at all sizes tried, up to 96 GiB
DMV	Successful commits up to 64 GiB, timeout at 5.5 h during 96 GiB trial

Figure 3: Wall-clock time to commit one large file to a fresh repository





Figure 4: Wall-clock time to commit many 1KiB files to a fresh repository



Number of Files

File Quantity Limits: inodes

Written: And we find that storing millions of objects loose as files with hash-based names will result in inefficient write speeds and use of disk space.

Git, Mercurial, DMV, and the copy operation all failed when trying to store 7.5 million files or more, reporting that the disk was full. However, the disk was not actually out of space — it was out of inodes.

Unix filesystems, ext4 included, store file and directory metadata in a data structure called an *inode*, which reside in a fixed-length table [19]. When all of the inodes in the table are allocated, the filesystem cannot store any more files or directories.

Bup avoided the inode limit. Bup trials could continue until the input data itself exhausted the system's inodes attempting to generate 25 million input files.

Otto:
Can be
configured
with more
inodes (max
in Linux
is...?)

Commit Times for Increasing Numbers of Files

Figure 4 shows the time required for the initial commit, storing all files into a fresh empty repository. Here we see the commit times for Git and DMV increasing quadratically with the number of files, while Mercurial, Bup, and the copy increase linearly.

5 Discussion

File Size Limits

Written: Discuss implementation details that lead to size limits

By using files as the basic unit of storage, and storing files as deltas against a base revision, both Git and Mercurial will at some point load an entire file into memory in order to compare it to another version. This limits the maximum file size that the system can work with to what can fit into RAM. In Mercurial's case, the error message that appears when attempting to commit a 2 GiB file warns that 6 GiB will be required to manage it. And because it has to calculate deltas in order to store a file at all, Mercurial

simply cannot work with any file that it can't fit into memory three times over. This is why Mercurial could not store files larger than 1.5 GiB in the file-size experiments (??).

Because Git's delta calculation happens behind-the-scenes in a secondary phase, it can still manage to commit files larger than available RAM, but it prints errors as the other operations fail. The two-phase approach also requires extra disk space and processing power. If a large file is changed, then both revisions will be written in full, taking twice the disk space. Then a separate operation will have to reread both blobs in full to calculate deltas and pack the objects.

Both DMV and Bup avoid these pitfalls by operating with a finer granularity, using a rolling hash to divide files into chunks by their content. It is the chunks and their indexes that must fit into memory, not the entire file. And then since chunks are only a few kilobytes and chunk indexes are hierarchical, file size becomes theoretically unlimited. Dividing into chunks by rolling hash also makes delta compression unnecessary, because identical chunks in different files or file revisions will naturally de-duplicate. At this point, it is the method of object storage that becomes the bottleneck.

DMV Sluggishness

DMV's parabolic increase is due to the way it breaks the large file into chunks and stores objects as individual files on the filesystem. While it is reading one large file, it is writing many small files, which incurs filesystem overhead. So its performance characteristic for storing a large file is closer to that of storing many files (??). Bup also breaks the file into many chunks, but it avoids the filesystem overhead by recombining the chunks into pack files. We investigate the filesystem overhead further in ??.

This sluggishness is due to the way DMV stores chunks of the file as individual files on the filesystem, turning the problem of storing one large file into the problem of storing many small files. Storing many small files in this way incurs filesystem overhead, as we discovered in the results of the number-of-files experiment (??), and later performed more experiments to examine in detail (??).

File Quantity Limits

Git, Mercurial, DMV, and the copy all create one file in their object stores for each input file. So to store 7.5 million files, they will create 7.5 million more, resulting in 15 million files on the filesystem, plus directories. However, the 197 GiB experiment partition has 13 107 200 total inodes, so storing 15 million files is impossible.

Bup is able to store more files because it does not write a separate object file for each input file. Bup aggregates its DAG objects into pack files, writing several large files instead many small files. As such, it does not exhaust the disk's inodes, and can continue until the experiment itself exhausts the system's inodes when it tries to go up from 10 million files to the next step and run a trial with 25 million files.

Hash-Based Directory Names Cause Disk Seeking

Written: Discuss Git and DMV slowing down due to random writes

We saw in the file-size commit times (??) that DMV's time increased quadratically, and we suspected that was because it was creating many small files and incurring filesystem overhead. This effect would explain why both Git and DMV do so poorly here while Bup would fare much better. But why then would Mercurial and the copy also have a linear increase instead of an quadratic one?

The difference is the naming schemes of stored files. Git and DMV name each object file according to the SHA-1 hash of the object's contents, while Mercurial, like the copy, uses the original input file's name. This means that Git and DMV write files in a random order with respect to their names, jumping between different object store subdirectories, while Mercurial and the copy can write files in the order they read them, one subdirectory at a time. The filesystem is most likely optimized for that kind of sequential write.

6 Related Work

Trim and compress the related work section

Distributed storage and synchronization systems

Write: [Related Works: Git](#)

Write: [Related Works: Mercurial](#)

Camlistore Camlistore [8] is an open-source project to create a private long-term data storage system for personal users. It allows storage of diverse types of data and it synchronizes between multiple replicas of the data store. However, it eschews normal filesystems and creates its own schemas to store various media.

Dat Data Dat [14] is an open-source project for publishing and sharing scientific data sets for research. This project has a lot of overlap with ours, and several of the core ideas are similar, including breaking files into smaller chunks, and tracking changes via a Git-like DAG. However, their focus is different. The Dat team is concentrating on publishing research data, and making that specific task as simple as possible for non-technical researchers who might not be familiar with version control. By contrast, our project operates at a lower level of abstraction, offering the full power of version control in a very general way, exposing and illuminating the complexities rather than trying to hide them or automate them away.

Where Dat focuses on publishing on the open internet, we focus on ad-hoc networks and data that may be private. Where Dat has components for automating peer discovery and consensus, we work at a lower level, trying to perfect and generalize the storage aspect first. Dat seems to assume that data sets will be small enough to fit on a typical disk on a workstation, while we want to scale even larger.

We hope that our system could be used as a base to build something like Dat, but we intend for DMV to be more general than the Dat core.

Eyo Eyo [22] is system for storing personal media and synchronizing it between devices. It utilizes a Git-like content-addressed object store behind the scenes, but it works more like a networked filesystem than version control. It focuses on organizing media by metadata, which requires agreement on metadata formats, and it requires applications to be rewritten to access files via Eyo rather than the filesystem, both of which are thorny and ambitious problems. We focus purely on storage and synchronization.

Git-Annex and Git-Media Git-annex [9] and Git-media [3] are open-source projects that extend Git with special handling for larger files. Both store information about the

larger files in the normal Git repository and then store the files themselves in a separate location. Git-media stores all the larger files in a separate data store which may be remote. Git-annex is more flexible. Annex files may be spread across several different remote repository clones or data stores, and Git-annex has features for tracking the locations of annex files in different remote repositories and moving them from one repository to another. These tracking and distribution features are very similar to our goals. However, Git-annex is not quite as flexible as we aim for with DMV. It considers the large files atomic units, and it does not break them into smaller chunks for de-duplication. Also, because metadata is processed by Git, it has the same limitations that Git does. All repositories must have all metadata, and performance suffers when metadata is too large to fit into RAM.

IPFS: The Interplanetary Filesystem IPFS [2] is an open-source project to create a global content-addressed filesystem. By its global nature, all files are stored publicly, in a global network of nodes with global addressing. IPFS should be an excellent resource for storing published information, but we want DMV to work with private data sets. We want individuals and organizations to be able manage their own data stores privately on their own hardware.

It should be noted that IPFS does have support for storing private objects by way of object-level encryption. However, this seems wasteful of disk space, since small changes in the plain text of a file would completely change the ciphertext, leaving no way to compress the redundancy.

Kademlia Kademlia [13] is an advanced distributed hash table system that updates its network topology information as part of normal lookups. Its focus is on efficient routing between nodes, while we are focusing first on storage.

Rsync Rsync [23] is a utility that synchronizes files across a low-bandwidth network link by using a rolling hash to find similar parts of the file, and then transfer only those portions of the file that have changed. Its rolling hash algorithm is the inspiration for the chunk-splitting algorithms used by Rsyncrypto [21], Bup [17], and DMV (??). As a utility, Rsync focuses exclusively on synchronizing files. It does not track history or the relation between files.

Rsync coauthor Andrew Tridgell is also indirectly responsible for the development of Git. His work to reverse-engineer the BitKeeper protocol is what caused BitMover to revoke the special free BitKeeper license for Linux kernel developers, which is what led Linus Torvalds to develop his own *source code manager (SCM)* [6].

Rsyncrypto Rsyncrypto [21] is an encryption system that uses a rolling hash to encrypt files by chunks, so that the encrypted files are still "rsync-able." Normally, a goal of an encryption algorithm is to have small changes in the input lead to a completely different ciphertext. This is good for security, but it negates the advantages of rsync — if the entire file is different, there are no advantages to be had in transferring only the changed parts when synchronizing files. Rsync uses a rolling hash to break the file into chunks and encrypt the chunks separately, so that a small change in the input will only change the ciphertext for the chunk that contains it, and so rsync can transfer only the chunk that changes.

Rsyncrypto's uses the same rolling hash algorithm as Rsync [21, 23], setting chunk boundaries where the checksum is zero. DMV uses this same chunking algorithm (see ??).

De-duplicating Storage and Backup

Boar Boar [7] is an open-source project to create a version control system for large binary files. It is one of the main inspirations for our project. It stores file versions in a content-addressed way, and provides de-duplication for large files that only change in small pieces, and it can truncate history to reclaim disk space. However, Boar retreats to the centralized version control paradigm, with a central repository that working directories must connect to in order to check files in or out. We want to provide the advantages of Boar in a flexible distributed version control model. Boar also has practical limitations on repository size and number of files. repositories are assumed to fit on one disk volume, and file metadata is assumed to fit into Ram. DMV tries to avoid both of those limitations.

Bup Bup [16] is an open-source file backup system that is based on Git's repository format. It does many of the things that DMV does: it breaks files into chunks by rolling hash, and it has considerations for metadata that is larger than RAM. Bup is one of the systems we evaluated along with Git and Mercurial (??), and we determined that Bup's manner of storing chunks is ideal for the next step for DMV (??). However, though Bup is built on version control, it is locked into a backup-based workflow. History is linear and based on clock time of backup. And it assumes that the whole data set and the whole repository can fit onto one filesystem. DMV is built to be more general and versatile.

Time Machine Time Machine [5] is a backup utility from Apple, included in Mac OS X Leopard, that makes frequent automated file hierarchy backups, using filesystem hard links to de-duplicate unchanged files [18]. This de-duplication allows it to store many different backup versions. Time Machine's functionality can be mimicked by using Rsync with the `--link-dest` option, which also hard-links unchanged files during a recursive sync [11].

7 Conclusion

Written: We conclude that the key to storing large files is to break them into smaller chunks, and that the key to storing many small chunks is to aggregate them into larger files.

We have performed experiments to probe the scalability limits of existing DAG-based distributed version control systems. We have shown that the maximum size of file that Git and Mercurial can store is limited by the amount of available memory in the system. We conclude that this is because those systems calculate deltas of files to de-duplicate data, and they load the entire file into memory in order to do so.

We have also rediscovered the limits of the Unix filesystem for storing many small files. We saw that writing files smaller than the filesystem block size incurs storage overhead, that splitting files among too many subdirectories takes inodes that are needed to store files, and that jumping between directories when writing files incurs write-speed penalties.

We have shown that any VCS that stores objects as individual files on the filesystem will encounter these filesystem limitations as they try to scale in terms of number of files. A VCS that also breaks files into chunks will turn the problem of storing large files into the problem of storing many files, again encountering these limitations. However, the limitations can be avoided by aggregating objects into pack files as Bup does.

Write: We intend to incorporate these insights into future versions of DMV.

References

- [1] BARUA, A., THOMAS, S. W., and HASSAN, A. E. “What are developers talking about? An analysis of topics and trends in Stack Overflow”. In: *Empirical Software Engineering* 19.3 (2014), pp. 619–654. ISSN: 1573-7616. DOI: 10.1007/s10664-012-9231-y. URL: <http://dx.doi.org/10.1007/s10664-012-9231-y>.
- [2] BENET, J. et al. *IPFS: The Interplanetary Filesystem*. GitHub. 2014. URL: <https://github.com/ipfs/ipfs>.
- [3] CHACON, S., LEBEDEV, A., et al. *git-media*. URL: <https://github.com/alebedev/git-media>.
- [4] CHACON, S. and STRAUB, B. *Pro Git*. 2nd. Berkely, CA, USA: Apress, 2014. ISBN: 1484200772, 9781484200773. URL: <https://git-scm.com/book/en/v2> (visited on Apr. 27, 2017).
- [5] CISLER, P. et al. *System for electronic backup*. US Patent App. 11/499,848. 2008. URL: <https://www.google.com/patents/US20080034004>.
- [6] CLOER, J. *10 Years of Git: An Interview with Git Createor Linus Torvalds*. Linux.com. Apr. 6, 2015. URL: <https://www.linux.com/blog/10-years-git-interview-git-creator-linus-torvalds> (visited on May 2, 2017).
- [7] EKBERG, M. et al. *Boar*. URL: <http://www.boarvcs.org/>.
- [8] FITZPATRICK, B. et al. *Camlistore is your personal storage system for life*. URL: <https://camlistore.org/>.
- [9] HESS, J. et al. *git-annex*. 2015. URL: <http://git-annex.branchable.com/>.
- [10] HUMBLE, J. and FARLEY, D. *Continuous Delivery: Reliable Software Releases through Build, Test, and Deployment Automation (Adobe Reader)*. Pearson Education, 2010.
- [11] JAKL, M. *Time Machine for every Unix out there*. Blog. Nov. 2007. URL: https://blog.interlinked.org/tutorials/rsync_time_machine.html (visited on May 12, 2017).
- [12] KARP, R. M. and RABIN, M. O. “Efficient randomized pattern-matching algorithms”. In: *IBM Journal of Research and Development* 31.2 (1987), pp. 249–260. ISSN: 0018-8646. DOI: 10.1147/rd.312.0249.
- [13] MAYMOUNKOV, P. and MAZIÈRES, D. “Kademlia: A Peer-to-Peer Information System Based on the XOR Metric”. In: *Peer-to-Peer Systems: First International-Workshop, IPTPS 2002 Cambridge, MA, USA, March 7–8, 2002 Revised Papers*. Ed. by DRUSCHEL, P., KAASHOEK, F., and ROWSTRON, A. Berlin, Heidelberg: Springer Berlin Heidelberg, 2002, pp. 53–65. ISBN: 978-3-540-45748-0. DOI: 10.1007/3-540-45748-8_5. URL: http://dx.doi.org/10.1007/3-540-45748-8_5.

- [14] OGDEN, M., BUUS, M., MCKELVEY, K., et al. *Dat Data*. URL: <http://dat-data.com/>.
- [15] O’SULLIVAN, B. *Mercurial: The Definitive Guide*. O’Reilly Media, Inc., 2009. ISBN: 0596800673, 9780596800673. URL: <http://hgbook.red-bean.com/>.
- [16] PENNARUN, A., BROWNING, R., et al. *Bup, it backs things up*. URL: <https://bup.github.io/> (visited on Apr. 26, 2017).
- [17] PENNARUN, A., BROWNING, R., et al. *The Crazy Hacker’s Crazy Guide to Bup Craziness*. “DESIGN” document in Bup source code. URL: <https://github.com/bup/bup/blob/master/DESIGN> (visited on Apr. 26, 2017).
- [18] POND, J. *How Time Machine Works its Magic*. Website. Aug. 31, 2013. URL: <http://pondini.org/TM/Works.html> (visited on May 12, 2017).
- [19] RITCHIE, O. M. and THOMPSON, K. “The UNIX time-sharing system”. In: *The Bell System Technical Journal* 57.6 (1978), pp. 1905–1929. ISSN: 0005-8580. DOI: 10.1002/j.1538-7305.1978.tb02136.x.
- [20] RUPARELIA, N. B. “The History of Version Control”. In: *SIGSOFT Softw. Eng. Notes* 35.1 (Jan. 2010), pp. 5–9. ISSN: 0163-5948. DOI: 10.1145/1668862.1668876. URL: <http://doi.acm.org/10.1145/1668862.1668876>.
- [21] SHEMESH, S. *Rsyncrypto algorithm*. Rsyncrypto home page. URL: <https://rsyncrypto.lingnu.com/index.php/Algorithm> (visited on Apr. 17, 2017).
- [22] STRAUSS, J. et al. “Eyo: Device-transparent Personal Storage”. In: *Proceedings of the 2011 USENIX Conference on USENIX Annual Technical Conference*. USENIXATC’11. Portland, OR: USENIX Association, 2011, pp. 35–35. URL: <http://dl.acm.org/citation.cfm?id=2002181.2002216>.
- [23] TRIDGELL, A. and MACKERRAS, P. *The rsync algorithm*. Tech. rep. Australian National University, 1996. URL: <http://hdl.handle.net/1885/40765> (visited on May 12, 2017).
- [24] WINSLOW, R. *Don’t ever commit binary files to Git! Or what to do if you do*. The Blog of Robin. June 11, 2013. URL: <https://robinwinslow.uk/2013/06/11/dont-ever-commit-binary-files-to-git/> (visited on May 11, 2017).

Find places where more citations can be added

Fix all broken links in text

Todo list

Incorporate abstract suggestion from Otto	1
Write: Distributed version control is an interesting form of distributed system because it takes eventual consistency to the extreme.	2
Write: Every replica of a repository contains the full history in an append-only data structure, any replica may add new commits, and conflicting updates are reconciled later in a merge operation.	2
Write: This paper explores the challenges of using version control to store larger binary files, with the goal of building a scalable, highly-available, distributed storage system for media files such as images, audio, and video.	2
Reduce and compress the DMV section	3

It's the rsync algorithm	4
Otto: Don't use "would" so much	5
Include platform table?	6
Otto: Don't use "would" so much	7
Otto: Give reason Mercurial needs 3x RAM	7
Ask Otto: Should that go in the discussion section?	7
Otto: Is this A, D, or I	7
Otto: Why timed out?	7
Otto: Make clearer that this was just experiment set up and can be changed	7
Otto: Why parabolic	7
Ask Otto: In discussion?	7
Remove this table?	8
Otto: Can be configured with more inodes (max in Linux is...?)	9
Trim and compress the related work section	11
Write: Related Works: Git	11
Write: Related Works: Mercurial	11
Write: We intend to incorporate these insights into future versions of DMV.	14
Find places where more citations can be added	15
Fix all broken links in text	15

Scratch Pad

This section contains text that I have written, but decided to set aside for now. It should disappear if the `final` option is applied to the document.